

Capstone -1 Fraud Analysis Using ML Algorithms

SUBMITTED BY: GERALDINE MATHEW

Problem Statement: In this capstone project, you are going to analyze customer's purchase data set and build a machine learning algorithm to detect fraud purchases. And also you are going to perform customer segmentation

Dataset: Purchase_Fraud_Data.csv

CAPSTONE PROJECT - 1 FRAUD ANALYSIS USING ML ALGORITHMS

Problem Statement: In this capstone project, you are going to analyze customer's purchase data set and build a machine learning algorithm to detect fraud purchases. And also you are going to perform customer segmentation

Dataset: Purchase_Fraud_Data.csv

Dataset Description: Every row in this dataset contains information about purchases from multiple customers. Along with purchase details, we also have customer basic details like gender, date of birth etc. Individual column description as follows

User_id: Customer unique id Signup_time: Date & Time at which the customer signup in the platform
Purchase_time: The latest purchase date & time from a customer Purchase_value: Total purchase amount
Device_id: Unique device ID from which purchase was done Source: Medium through which customers reached the platform Browser: Browser used while purchasing IP_address: IP Address from which purchase was done Class: 1 = Target class; Fraud transaction; 0=Regular transaction Category: Type of product purchased Dob: Date of birth of the customer

Exploratory Data Analysis

1. Summarize numerical, categorical and date columns separately and list down your inferences
2. Identify and perform missing value treatment using appropriate techniques
3. Univariate analysis: For each column perform appropriate univariate analysis. (i.e. perform distribution analysis on numerical columns and frequency analysis on categorical columns)
4. Multivariate analysis: Take combinations of multiple columns and identify the relationship between them. a. Categorical vs numerical columns – bar charts, boxplots b. Numerical vs numerical columns – scatter plot c. Categorical vs multiple numerical columns – scatter plot d. Correlation matrix e. ...
5. Perform statistical hypothesis test to identify the relationship between input and target variables

Base Model for Benchmark

1. Convert all the input columns to numerical columns using one hot encoding
2. Make sure you temporarily ignore the date columns and id columns
3. Split the data in training and testing
4. Build a simple decision tree with max_depth=5 and identify accuracy and f1 score. Keep this as a bench mark values to explain the improvement post feature engineering

Feature Engineering

Using the original data, perform the following feature engineering techniques to create new columns for modelling

1. Using all the date columns, extract the following information wherever necessary a. Year, month, day, hour, day of the week (mon, tue, etc)
2. Using date of birth column, try to calculate the appropriate age of the customer. (Compute difference between purchase date and date of birth)
3. Using age column, identify age buckets using binning method
4. Compute no. of hours/days between purchase time and signup time

5. For all appropriate numerical variables, compute buckets using binning method

Model Building

1. Convert all the input variables to numeric
2. Split the new data in to training and testing
3. Build the following classification models using new data a. Decision Trees b. Random Forest c. Boosting techniques d. XGBoost
4. For all the above classifiers, make sure that you perform hyper parameter tuning to select optimal value hyper parameters
5. For all the above methods, report accuracy and f1 score
6. Also in a single plot, compare the ROC curves for all the above models
7. Choose an appropriate model based on above inferences and report the same

Customer Segmentation

The company would like to segment their customers using various attributes, so that they can perform targeted marketing campaign. Do the following to group the customers in the different clusters

1. In the original feature engineered dataset, exclude the following columns: ids, dates, class (target column)
2. Convert all the input variables to numeric
3. Using Elbow method identify optimal number of clusters required to group customers in to various clusters
4. Perform K-Means clustering using appropriate number of clusters
5. Tag each customer with a cluster
6. Cluster-wise report average values of all input variables

```
In [1]: import numpy as np
import pandas as pd
from IPython.display import display, HTML
%matplotlib inline
import datetime
import seaborn as sns
from sklearn.model_selection import cross_val_score
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, precision_score
```

In []:

In []: *#reading the file*

```
In [2]: df = pd.read_csv('Purchase_Fraud_data.csv')
df.head(5)
```

Out[2]:

	user_id	signup_time	purchase_time	purchase_value	device_id	source	browser	sex	ip_address	class
0	22058	2015-02-24 22:55:49	2015-04-18 02:47:11	65278	QVPSPJUOCKZAR	SEO	Chrome	M	7.327584e+08	
1	333320	2015-06-07 20:39:50	2015-06-08 01:38:54	96399	EOGFQPIZPYXFZ	Ads	Chrome	F	3.503114e+08	
2	1359	2015-01-01 18:52:44	2015-01-01 18:52:45	57296	YSSKYOSJHPPLJ	SEO	Opera	M	2.621474e+09	
3	150084	2015-04-28 21:13:25	2015-05-04 13:54:50	43650	ATGTXYKYKUDUQN	SEO	Safari	M	3.840542e+09	
4	221365	2015-07-21 07:09:52	2015-09-09 18:40:53	45016	NAUITBZFJKHWW	Ads	Safari	M	4.155831e+08	

```
In [3]: df1=df.copy()
```

```
In [4]: #Summarize numerical, categorical and date columns separately and list down your inferences
```

```
In [5]: df.shape
```

```
Out[5]: (151112, 12)
```

```
In [6]: df.count()
```

```
Out[6]: user_id      151112
signup_time    151112
purchase_time   151112
purchase_value  151112
device_id      151112
source         151112
browser        151112
sex            146185
ip_address     151112
class          151112
category       151112
dob            146188
dtype: int64
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	user_id	purchase_value	ip_address	class
count	151112.000000	151112.000000	1.511120e+05	151112.000000
mean	200171.040970	50521.469003	2.152145e+09	0.093646
std	115369.285024	28533.667117	1.248497e+09	0.291336
min	2.000000	1016.000000	5.209350e+04	0.000000
25%	100642.500000	25919.000000	1.085934e+09	0.000000
50%	199958.000000	50484.000000	2.154770e+09	0.000000
75%	300054.000000	75296.250000	3.243258e+09	0.000000
max	400000.000000	100092.000000	4.294850e+09	1.000000

```
In [8]: #Identify and perform missing value treatment using appropriate techniques
```

```
In [9]: df.isna().sum()
```

```
Out[9]: user_id      0
signup_time    0
purchase_time   0
purchase_value  0
device_id      0
source         0
browser        0
sex            4927
ip_address     0
class          0
category       0
dob            4924
dtype: int64
```

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 151112 entries, 0 to 151111
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   user_id                151112 non-null  int64
1   signup_time            151112 non-null  object
2   purchase_time          151112 non-null  object
3   purchase_value         151112 non-null  int64
4   device_id              151112 non-null  object
5   source                 151112 non-null  object
6   browser                151112 non-null  object
7   sex                    146185 non-null  object
8   ip_address             151112 non-null  float64
9   class                  151112 non-null  int64
10  category               151112 non-null  object
11  dob                    146188 non-null  object
dtypes: float64(1), int64(3), object(8)
memory usage: 13.8+ MB
```

```
In [11]: df.isna().sum()/len(df)*100
#we can see the null values in Sex and Dob columns
```

```
Out[11]: user_id                0.000000
signup_time            0.000000
purchase_time          0.000000
purchase_value         0.000000
device_id              0.000000
source                 0.000000
browser                0.000000
sex                    3.260496
ip_address             0.000000
class                  0.000000
category               0.000000
dob                    3.258510
dtype: float64
```

```
In [12]: #Since it has null values less than 10 % we can drop null values of sex and dob columns.
```


```
In [13]: df['class'].value_counts()
```

```
Out[13]: 0    136961
1     14151
Name: class, dtype: int64
```

```
In [14]: df = df.dropna()
df.head()
```

```
Out[14]:
```

	user_id	signup_time	purchase_time	purchase_value	device_id	source	browser	sex	ip_address	class
0	22058	2015-02-24 22:55:49	2015-04-18 02:47:11	65278	QVPSPJUOCKZAR	SEO	Chrome	M	7.327584e+08	
1	333320	2015-06-07 20:39:50	2015-06-08 01:38:54	96399	EOGFQPIZPYXFZ	Ads	Chrome	F	3.503114e+08	
2	1359	2015-01-01 18:52:44	2015-01-01 18:52:45	57296	YSSKYOSJHPPLJ	SEO	Opera	M	2.621474e+09	
3	150084	2015-04-28 21:13:25	2015-05-04 13:54:50	43650	ATGTXYKYKUDUQN	SEO	Safari	M	3.840542e+09	
4	221365	2015-07-21 07:09:52	2015-09-09 18:40:53	45016	NAUITBZFJKHWW	Ads	Safari	M	4.155831e+08	



```
In [15]: df.columns
```

```
Out[15]: Index(['user_id', 'signup_time', 'purchase_time', 'purchase_value',
               'device_id', 'source', 'browser', 'sex', 'ip_address', 'class',
               'category', 'dob'],
              dtype='object')
```

```
In [16]: df['class'].value_counts()
```

```
Out[16]: 0    128164
         1     13240
         Name: class, dtype: int64
```

```
In [17]: df.isna().sum()/len(df)*100
```

```
Out[17]: user_id          0.0
         signup_time      0.0
         purchase_time     0.0
         purchase_value    0.0
         device_id        0.0
         source           0.0
         browser          0.0
         sex              0.0
         ip_address       0.0
         class            0.0
         category         0.0
         dob              0.0
         dtype: float64
```

```
In [18]: df.skew()
```

C:\Users\asus\AppData\Local\Temp\ipykernel_13324\1665899112.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
df.skew()
```

```
Out[18]: user_id          0.001247
         purchase_value    0.000862
         ip_address       -0.007756
         class            2.789898
         dtype: float64
```

```
In [19]: df.std()
```

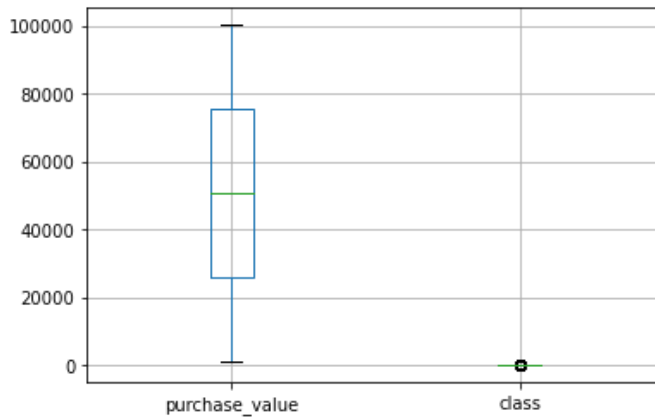
C:\Users\asus\AppData\Local\Temp\ipykernel_13324\3390915376.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
df.std()

```
Out[19]: user_id          1.153234e+05  
purchase_value  2.853566e+04  
ip_address      1.248164e+09  
class           2.913177e-01  
dtype: float64
```

Outlier visualization

```
In [20]: df.drop(['user_id', 'ip_address'], axis=1).boxplot()  
# we can observe that no outlier present.
```

```
Out[20]: <AxesSubplot:>
```



```
In [21]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 141404 entries, 0 to 151111  
Data columns (total 12 columns):  
#   Column          Non-Null Count  Dtype  
---  ---          -  
0   user_id         141404 non-null  int64  
1   signup_time     141404 non-null  object  
2   purchase_time   141404 non-null  object  
3   purchase_value  141404 non-null  int64  
4   device_id       141404 non-null  object  
5   source          141404 non-null  object  
6   browser         141404 non-null  object  
7   sex             141404 non-null  object  
8   ip_address      141404 non-null  float64  
9   class           141404 non-null  int64  
10  category        141404 non-null  object  
11  dob             141404 non-null  object  
dtypes: float64(1), int64(3), object(8)  
memory usage: 14.0+ MB
```

```
In [22]: obj=df.select_dtypes(include='object').columns
obj
```

```
Out[22]: Index(['signup_time', 'purchase_time', 'device_id', 'source', 'browser', 'sex',
               'category', 'dob'],
              dtype='object')
```


```
In [23]: df1['class']=df1['class'].astype('object')
```

UNIVARIATE ANALYSIS OF NUMERICAL COLUMN

```
In [24]: num=['purchase_value']
```

```
In [25]: a = 4
b = 2
counter = 1
plt.figure(figsize = [12, 8])


for i in num:
    plt.subplot(a,b,counter)
    plt.title(i)
    sns.kdeplot(df1.loc[:, i])
    counter = counter+1
plt.tight_layout()
plt.show()
```




```
In [26]: # Purchase value is highly skewed
```

UNIVARIATE ANALYSIS OF CATEGORICAL COLUMN

```
In [27]: x=df1['source'].value_counts()
plt.pie(x,labels=x.index,autopct='%0.2f',shadow=True)
plt.show()
```




```
In [28]: # SEO is highest in percentage followed by Ads and Direct.
x=df1['sex'].value_counts()
plt.pie(x,labels=x.index,autopct='%0.2f',shadow=True)
plt.show()
```




```
In [29]: # In terms of gender the category of Males are more in percentage as compared to females.
```

```
In [30]: x=df1['browser'].value_counts()
plt.pie(x,labels=x.index,autopct='%0.2f',shadow=True)
plt.show()
```



```
In [31]: #majority of the customers are using Chrome followed by IE,Safari,Firefox and Opera.
```

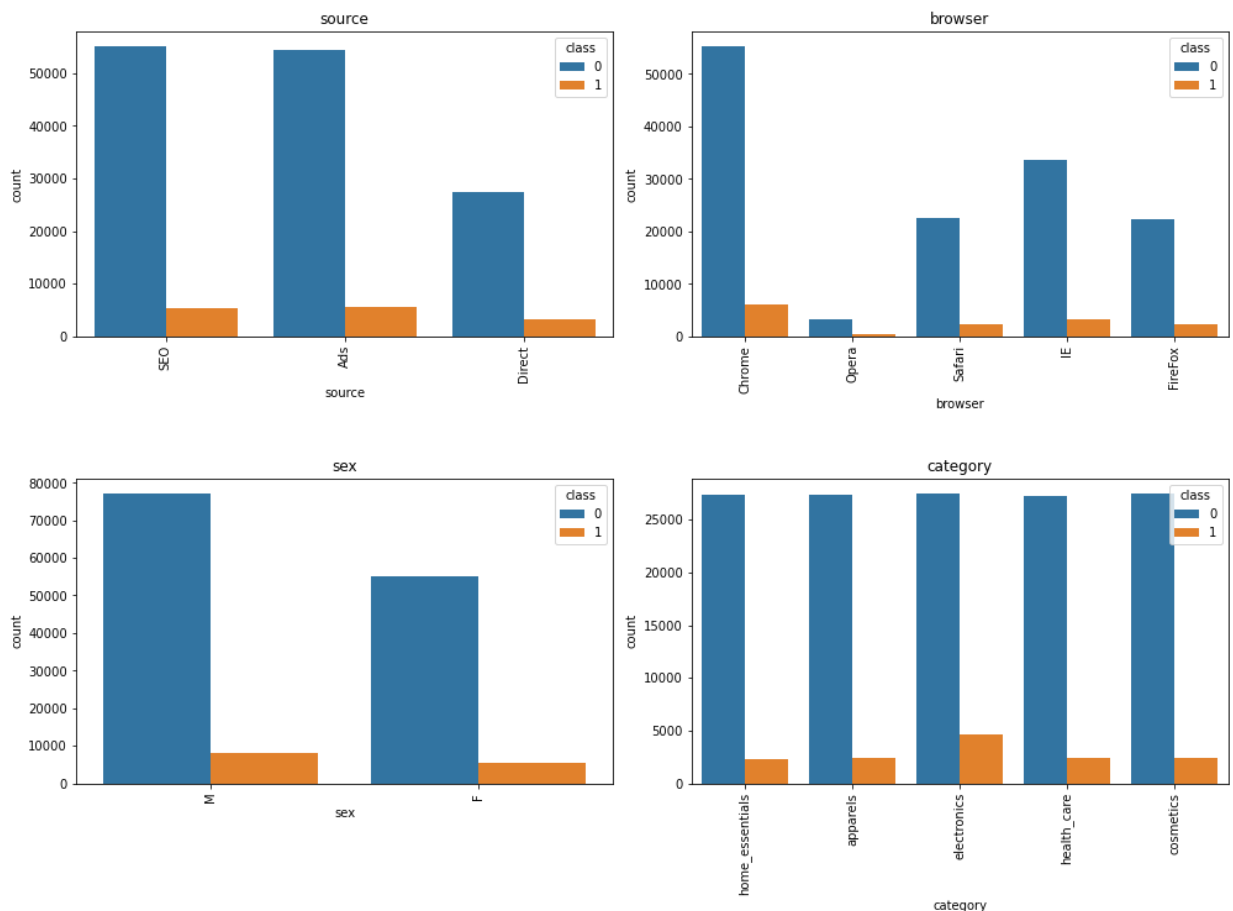
```
In [32]: x=df1['category'].value_counts()
plt.pie(x,labels=x.index,autopct='%0.2f',shadow=True)
plt.show()
```



Bi variate analysis of the categorical columns.Customer segmentation on the basis of source,browser,sex,category

```
In [33]: obj=[ 'source', 'browser', 'sex','category']
```

```
In [34]: cols=obj
a = 3
b = 2
counter = 1
plt.figure(figsize = [14, 14])
for i in cols:
    plt.subplot(a,b,counter)
    plt.title(i)
    sns.countplot(x=df1[i],hue=df1['class'])
    counter = counter+1
    plt.xticks(rotation=90)
plt.tight_layout()
plt.show()
```



Bi variate analysis of the continuous columns. Customer segmentation on the basis of purchase value


```
In [35]: num=df1.select_dtypes(exclude='object').columns
```

```
In [36]: num[1]
```

```
Out[36]: 'purchase_value'
```

```
In [37]: cols=num[1]
plt.figure(figsize=[8,6])
plt.title(i)
sns.barplot(y=df1['purchase_value'],x=df1['class'])
```

Out[37]: <AxesSubplot:title={ 'center': 'category'}, xlabel='class', ylabel='purchase_value'>



In [38]: *#from the above graph we can infer that approximately equal amount of Average purcahse is done*

```
In [39]: df1['purchase_value'].min()
```

Out[39]: 1016

```
In [40]: df1['purchase_value'].median()
```

Out[40]: 50484.0

```
In [41]: df1['purchase_value'].mean()
```

Out[41]: 50521.46900312351





```
In [42]: df1['purchase_value'].max()
```

Out[42]: 100092

```
In [43]: import matplotlib.pyplot as plt
%matplotlib inline


df.hist(figsize=(20,30))
```

```
Out[43]: array([[<AxesSubplot:title={'center':'user_id'}>,
<AxesSubplot:title={'center':'purchase_value'}>],
[<AxesSubplot:title={'center':'ip_address'}>,
<AxesSubplot:title={'center':'class'}>]], dtype=object)
```



```
In [44]: plt.figure(figsize=[8,8])
sns.heatmap(df.corr()[df.corr().>0.5],annot=True)
```

Out[44]: <AxesSubplot:>



No strong co-relation with the target class

CREATING BASE MODEL FOR BENCH MARK

```
In [45]: df['source'] = df['source'].replace({'SEO':0, 'Direct':1, 'Ads':3})
```

```
In [46]: df['browser'] = df['browser'].replace({'Chrome':0, 'Opera':1, 'Safari':3, 'IE':4, 'FireFox':5})
```

```
In [47]: df['sex'] = df['sex'].replace({'M':0, 'F':1})
```

```
In [48]: df['category'] = df['category'].replace({'home_essentials':0, 'apparels':1, 'electronics':3, 'he
```

```
In [49]: df1=df.drop(['user_id', 'signup_time', 'device_id', 'purchase_time',
                    'ip_address',
                    'dob'],axis=1)
```

```
In [50]: import pandas as pd
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix, auc, roc_curve
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
```

```
In [51]: df1
```

```
Out[51]:
```

	purchase_value	source	browser	sex	class	category
0	65278	0	0	0	0	0
1	96399	3	0	1	0	1
2	57296	0	1	0	1	3
3	43650	0	3	0	0	4
4	45016	3	3	0	0	0
...
151107	58549	0	0	0	1	1
151108	57952	0	3	0	0	3
151109	19003	0	4	1	0	4
151110	68296	0	0	0	0	3
151111	23622	1	4	0	0	1

141404 rows × 6 columns

```
In [52]: target_col = 'class'
input_cols = df1.drop(target_col, axis=1).columns
train_x, test_x, train_y, test_y = train_test_split(df1[input_cols],
                                                    df1[target_col],
                                                    test_size=0.2,
                                                    random_state=1)

train_x.shape, test_x.shape, train_y.shape, test_y.shape
```

```
Out[52]: ((113123, 5), (28281, 5), (113123,), (28281,))
```

```
In [53]: scaler = StandardScaler().fit(train_x)
train_x_scaled = pd.DataFrame(scaler.transform(train_x),
                              columns=train_x.columns,
                              index=train_x.index)
test_x_scaled = pd.DataFrame(scaler.transform(test_x),
                             columns=test_x.columns,
                             index=test_x.index)
```

```
In [54]: def get_model_metrics(actual, predicted, classifier_name):
    acc = accuracy_score(actual, predicted)
    f1 = f1_score(actual, predicted)
    tn, fp, fn, tp = confusion_matrix(actual, predicted).ravel()
    sensitivity = tp / (tp + fn)
    specificity = tn / (tn + fp)
    fpr, tpr, thresholds = roc_curve(actual, predicted, pos_label=1)
    auc_value = auc(fpr, tpr)
    return {'accuracy': acc,
            'f1-score': f1,
            'classifier': classifier_name,
            'sensitivity': sensitivity,
            'specificity': specificity,
            'AUC': auc_value,
            'fpr_values': fpr,
            'tpr_values': tpr}
```

```
In [55]: res_cols = ['classifier', 'f1-score', 'sensitivity', 'specificity', 'AUC', 'accuracy']
df_results = pd.DataFrame(columns=res_cols)
```

```
In [56]: #USING SIMPLE DECISION TREE WITH MAX_DEPTH=5
```

```
In [57]: # Decision Trees
classifier_name = 'Decision Tree'
model = DecisionTreeClassifier(max_depth=5, random_state=10).fit(train_x_scaled, train_y)
test_y_pred = model.predict(test_x_scaled)
dt_metrics = get_model_metrics(test_y, test_y_pred, classifier_name)
print(dt_metrics)
if classifier_name not in df_results['classifier'].values:
    df_results = df_results.append(dt_metrics, ignore_index=True)

{'accuracy': 0.9066157490894947, 'f1-score': 0.0, 'classifier': 'Decision Tree', 'sensitivity': 0.0, 'specificity': 0.999960999961, 'AUC': 0.4999804999805, 'fpr_values': array([0.000000e+00, 3.9000039e-05, 1.0000000e+00]), 'tpr_values': array([0., 0., 1.])}
```


```
In [58]: df_results.drop(['fpr_values', 'tpr_values'], axis=1)
```

Out[58]:

	classifier	f1-score	sensitivity	specificity	AUC	accuracy
0	Decision Tree	0.0	0.0	0.999961	0.49998	0.906616

```
In [59]: plt.plot(dt_metrics['fpr_values'], dt_metrics['tpr_values'])
plt.legend(['Decision Tree; AUC: %.2f' % dt_metrics['AUC']])
```

Out[59]: <matplotlib.legend.Legend at 0x24e6855fe80>



In []:

PERFORM FEATURE ENGINEERING

Using the original data, perform the following feature engineering techniques to create new columns for modelling

Using all the date columns, extract the following information wherever necessary a. Year, month, day, hour, day of the week (mon, tue, etc) Using date of birth column, try to calculate the appropriate age of the customer. (Compute difference between purchase date and date of birth) Using age column, identify age buckets using binning method Compute no. of hours/days between purchase time and signup time For all appropriate numerical variables, compute buckets using binning method

```
In [60]: df2 = pd.read_csv('Purchase_Fraud_data.csv')
df2.head(5)
```

Out[60]:

	user_id	signup_time	purchase_time	purchase_value	device_id	source	browser	sex	ip_address	class
0	22058	2015-02-24 22:55:49	2015-04-18 02:47:11	65278	QVPSPJUOCKZAR	SEO	Chrome	M	7.327584e+08	
1	333320	2015-06-07 20:39:50	2015-06-08 01:38:54	96399	EOGFQPIZPYXFZ	Ads	Chrome	F	3.503114e+08	
2	1359	2015-01-01 18:52:44	2015-01-01 18:52:45	57296	YSSKYOSJHPPLJ	SEO	Opera	M	2.621474e+09	
3	150084	2015-04-28 21:13:25	2015-05-04 13:54:50	43650	ATGTXKYKUDUQN	SEO	Safari	M	3.840542e+09	
4	221365	2015-07-21 07:09:52	2015-09-09 18:40:53	45016	NAUITBZFJKHWW	Ads	Safari	M	4.155831e+08	

```
In [61]: df2.isna().sum()
```

```
Out[61]: user_id          0
signup_time         0
purchase_time        0
purchase_value       0
device_id           0
source              0
browser             0
sex                 4927
ip_address          0
class              0
category            0
dob                 4924
dtype: int64
```



```
In [62]: df2.isna().sum()/len(df2)*100
#we can see the null values in Sex and Dob
```

```
Out[62]: user_id          0.000000
signup_time      0.000000
purchase_time    0.000000
purchase_value   0.000000
device_id        0.000000
source           0.000000
browser          0.000000
sex              3.260496
ip_address       0.000000
class            0.000000
category         0.000000
dob              3.258510
dtype: float64
```

```
In [63]: df2= df2.dropna()
df2.head()
```


```
Out[63]:
```

	user_id	signup_time	purchase_time	purchase_value	device_id	source	browser	sex	ip_address	class
0	22058	2015-02-24 22:55:49	2015-04-18 02:47:11	65278	QVPSPJUOCKZAR	SEO	Chrome	M	7.327584e+08	
1	333320	2015-06-07 20:39:50	2015-06-08 01:38:54	96399	EOGFQPIZPYXFZ	Ads	Chrome	F	3.503114e+08	
2	1359	2015-01-01 18:52:44	2015-01-01 18:52:45	57296	YSSKYOSJHPPLJ	SEO	Opera	M	2.621474e+09	
3	150084	2015-04-28 21:13:25	2015-05-04 13:54:50	43650	ATGTXXKYKUDUQN	SEO	Safari	M	3.840542e+09	
4	221365	2015-07-21 07:09:52	2015-09-09 18:40:53	45016	NAUITBZFJKHWW	Ads	Safari	M	4.155831e+08	

```
In [64]: #Binning of Purchase Column
```

```
In [65]: df2['Purchase_Amount_Cat']=pd.cut(df2['purchase_value'],bins=[1016 ,20000,50000,100092],
labels=['purchase_value in 1k and 20k','purchase_value in 20k
```

```
In [66]: (pd.crosstab(df2['Purchase_Amount_Cat'],df2['class'],normalize=True)*100).plot(kind="bar")
plt.show()
```



```
In [67]: # from the above graph we can infer that Most of the regular transactions are taking place be
# Most Fraud Transaction are happening in range 50 k and 100 k
```

```
In [ ]:
```

```
In [68]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 141404 entries, 0 to 151111
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   user_id                141404 non-null int64
1   signup_time            141404 non-null object
2   purchase_time          141404 non-null object
3   purchase_value         141404 non-null int64
4   device_id              141404 non-null object
5   source                 141404 non-null object
6   browser                141404 non-null object
7   sex                   141404 non-null object
8   ip_address             141404 non-null float64
9   class                  141404 non-null int64
10  category                141404 non-null object
11  dob                    141404 non-null object
12  Purchase_Amount_Cat    141403 non-null category
dtypes: category(1), float64(1), int64(3), object(8)
memory usage: 14.2+ MB
```

```
In [69]: df2.head()
```

Out[69]:

	user_id	signup_time	purchase_time	purchase_value	device_id	source	browser	sex	ip_address	clk
0	22058	2015-02-24 22:55:49	2015-04-18 02:47:11	65278	QVPSPJUOCKZAR	SEO	Chrome	M	7.327584e+08	
1	333320	2015-06-07 20:39:50	2015-06-08 01:38:54	96399	EOGFQIPZYXFZ	Ads	Chrome	F	3.503114e+08	
2	1359	2015-01-01 18:52:44	2015-01-01 18:52:45	57296	YSSKYOSJHPPLJ	SEO	Opera	M	2.621474e+09	
3	150084	2015-04-28 21:13:25	2015-05-04 13:54:50	43650	ATGTXYKYKUDUQN	SEO	Safari	M	3.840542e+09	
4	221365	2015-07-21 07:09:52	2015-09-09 18:40:53	45016	NAUITBZFJKHWW	Ads	Safari	M	4.155831e+08	

```
In [70]: df2['class'].value_counts()
```

Out[70]:

```
0    128164
1     13240
Name: class, dtype: int64
```

```
In [71]: #We can see that our target class is imbalanced
```

```
In [ ]:
```

Conversion Of Date Time of Signup and Purchase Time

```
In [72]: import datetime as dt
df2['tot_mins_diff']=pd.to_datetime(df2['purchase_time'])-pd.to_datetime(df2['signup_time'])

df2.head()

df2['tot_mins_diff'] = df2['tot_mins_diff'].dt.total_seconds()

# Conversion of Difference in Minutes

df2['tot_mins_diff']=df2['tot_mins_diff']/60

df2.head()
```

Out[72]:


	user_id	signup_time	purchase_time	purchase_value	device_id	source	browser	sex	ip_address	clk
0	22058	2015-02-24 22:55:49	2015-04-18 02:47:11	65278	QVPSPJUOCKZAR	SEO	Chrome	M	7.327584e+08	
1	333320	2015-06-07 20:39:50	2015-06-08 01:38:54	96399	EOGFQIPZYXFZ	Ads	Chrome	F	3.503114e+08	
2	1359	2015-01-01 18:52:44	2015-01-01 18:52:45	57296	YSSKYOSJHPPLJ	SEO	Opera	M	2.621474e+09	
3	150084	2015-04-28 21:13:25	2015-05-04 13:54:50	43650	ATGTXYKYKUDUQN	SEO	Safari	M	3.840542e+09	
4	221365	2015-07-21 07:09:52	2015-09-09 18:40:53	45016	NAUITBZFJKHWW	Ads	Safari	M	4.155831e+08	

In [73]: df2

Out[73]:

	user_id	signup_time	purchase_time	purchase_value	device_id	source	browser	sex	ip_addres
0	22058	2015-02-24 22:55:49	2015-04-18 02:47:11	65278	QVPSPJUOCKZAR	SEO	Chrome	M	7.327584e+0
1	333320	2015-06-07 20:39:50	2015-06-08 01:38:54	96399	EOGFQPIZPYXFZ	Ads	Chrome	F	3.503114e+0
2	1359	2015-01-01 18:52:44	2015-01-01 18:52:45	57296	YSSKYOSJHPPLJ	SEO	Opera	M	2.621474e+0
3	150084	2015-04-28 21:13:25	2015-05-04 13:54:50	43650	ATGTXYKYKUDUQN	SEO	Safari	M	3.840542e+0
4	221365	2015-07-21 07:09:52	2015-09-09 18:40:53	45016	NAUITBZFJKHWW	Ads	Safari	M	4.155831e+0
...
151107	345170	2015-01-27 03:03:34	2015-03-29 00:30:47	58549	XPSKTWGPWINLR	SEO	Chrome	M	3.451155e+0
151108	274471	2015-05-15 17:43:29	2015-05-26 12:24:39	57952	LYSFABUCPCGBA	SEO	Safari	M	2.439047e+0
151109	368416	2015-03-03 23:07:31	2015-05-20 07:07:47	19003	MEQHCSJUBRBF	SEO	IE	F	2.748471e+0
151110	207709	2015-07-09 20:06:07	2015-09-07 09:34:46	68296	CMCXFGRHYSTVJ	SEO	Chrome	M	3.601175e+0
151111	138208	2015-06-10 07:02:20	2015-07-21 02:03:53	23622	ZINIADFCLHYPG	Direct	IE	M	4.103825e+0

141404 rows × 14 columns



In [74]: `#df2.drop(['signup_time', 'purchase_time'],axis=1,inplace=True)`

Customer segmentation on the basis of difference between signup time and purchase time

In [75]: `df2['tot_mins_diff'].max()`

Out[75]: 172799.53333333333

In [76]: `df2['tot_mins_diff'].min()`

Out[76]: 0.016666666666666666

In [77]: `df2['tot_mins_diff'].mean()`


Out[77]: 81064.4638903388

In [78]: `df2['tot_mins_diff'].median()`

Out[78]: 80805.4

```
In [79]: df2['total_mins_diff']=pd.cut(df2['tot_mins_diff'],bins=[0.01,1000,50000,173000],
                                         labels=['time between 0 to 1000 min','time between 1000 to 500
```

```
In [80]: (pd.crosstab(df2['total_mins_diff'],df2['class'],normalize=True)*100).plot(kind="bar")
plt.show()
```



Most of the fraud transactions are between 0 to 1000 mins of interval.
We can also observe that fraud people are taking less time than people who are making regular transaction.
people who are making Regular transactions are taking more time than 1000 mins.

Customer segmentation on the basis of different age group

```
In [82]: import datetime as dt
df2['dob'] = pd.to_datetime(df2['dob'], format = "%d-%m-%Y")
now = pd.Timestamp('now')
df2['age'] = (now - df2['dob'])
df2['age'] = df2['age'].astype(str)
df2[['age', 'age_waste']] = df2['age'].str.split("days", expand=True)
```

```
In [83]: df2 = df2.drop(['age_waste'],axis = 1)
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 141404 entries, 0 to 151111
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   user_id                141404 non-null  int64
1   signup_time            141404 non-null  object
2   purchase_time          141404 non-null  object
3   purchase_value         141404 non-null  int64
4   device_id              141404 non-null  object
5   source                 141404 non-null  object
6   browser                141404 non-null  object
7   sex                    141404 non-null  object
8   ip_address             141404 non-null  float64
9   class                  141404 non-null  int64
10  category               141404 non-null  object
11  dob                    141404 non-null  datetime64[ns]
12  Purchase_Amount_Cat    141403 non-null  category
13  tot_mins_diff          141404 non-null  float64
14  total_mins_diff        141404 non-null  category
15  age                    141404 non-null  object
dtypes: category(2), datetime64[ns](1), float64(2), int64(3), object(8)
memory usage: 16.5+ MB
```

```
In [84]: df2['age'] = pd.to_numeric(df2['age'], errors='coerce')
```

```
In [85]: df2['age']=df2['age']/365
```

```
In [86]: df2['age'].max()
```

```
Out[86]: 84.06027397260274
```

```
In [87]: df2['age'].min()
```

```
Out[87]: 25.405479452054795
```

```
In [88]: df2['age'].mean()
```


```
Out[88]: 41.74240983669481
```

```
In [89]: df2['age'].median()
```

```
Out[89]: 40.772602739726025
```

```
In [90]: df2['age']=pd.cut(df2['age'],bins=[20,35,50,83],
                        labels=['age between 20 to 35','age between 35 to 50','age bet
```

```
In [91]: (pd.crosstab(df2['age'],df2['class'],normalize=True)*100).plot(kind="bar")
plt.show()
```



Fraud transactions are higher between the ages of 50 to 83.

Most of the regular transactions are done by the age group of 35 to 50.

```
In [ ]:
```

```
In [92]: df2.drop(['signup_time','purchase_time', 'dob','device_id','ip_address','user_id','total_mins
```


```
In [93]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 141404 entries, 0 to 151111
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   purchase_value  141404 non-null  int64
1   source          141404 non-null  object
2   browser         141404 non-null  object
3   sex             141404 non-null  object
4   class           141404 non-null  int64
5   category        141404 non-null  object
6   tot_mins_diff   141404 non-null  float64
dtypes: float64(1), int64(2), object(4)
memory usage: 8.6+ MB
```

```
In [94]: # Here We are converting categorical data into dummy or indicator variables.
df2=pd.get_dummies(df2,drop_first=True)
```

```
In [95]: plt.figure(figsize=[8,8])
sns.heatmap(df2.corr()[df2.corr().>0.5],annot=True)
```

Out[95]: <AxesSubplot:>



```
In [96]: # No strong co-relation with the target class
```

#Perform statistical hypothesis test to identify the relationship between input and target variables

Plotting Histogram for every columns


```
In [97]: import matplotlib.pyplot as plt
        %matplotlib inline

        df2.hist(figsize=(20,30))
```

```
Out[97]: array([[<AxesSubplot:title={'center': 'purchase_value'}>,
                <AxesSubplot:title={'center': 'class'}>,
                <AxesSubplot:title={'center': 'tot_mins_diff'}>,
                <AxesSubplot:title={'center': 'source_Direct'}>],
               [<AxesSubplot:title={'center': 'source_SEO'}>,
                <AxesSubplot:title={'center': 'browser_FireFox'}>,
                <AxesSubplot:title={'center': 'browser_IE'}>,
                <AxesSubplot:title={'center': 'browser_Opera'}>],
               [<AxesSubplot:title={'center': 'browser_Safari'}>,
                <AxesSubplot:title={'center': 'sex_M'}>,
                <AxesSubplot:title={'center': 'category_cosmetics'}>,
                <AxesSubplot:title={'center': 'category_electronics'}>],
               [<AxesSubplot:title={'center': 'category_health_care'}>,
                <AxesSubplot:title={'center': 'category_home_essentials'}>,
                <AxesSubplot:~>, <AxesSubplot:~>]], dtype=object)
```



In []:

In []:

MODEL AFTER PERFORMING FEATURE ENGINEERING

Model Building

Convert all the input variables to numeric Split the new data in to training and testing Build the following classification models using new data a. Decision Trees b. Random Forest c. Boosting techniques d. XGBoost For all the above classifiers, make sure that you perform hyper parameter tuning to select optimal value hyper parameters For all the above methods, report accuracy and f1 score Also in a single plot, compare the ROC curves for all the above models Choose an appropriate model based on above inferences and report the same

```
In [98]: from sklearn.model_selection import train_test_split
```

```
In [99]: X=df2.drop('class',axis=1)
        y=df2['class']
```

```
In [100]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=10)
```

```
In [101]: from sklearn.metrics import roc_curve, auc
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import roc_auc_score
        from sklearn.linear_model import LogisticRegression
```

```
In [102]: df2['class'].value_counts()
```

```
Out[102]: 0    128164
         1    13240
         Name: class, dtype: int64
```

```
In [103]: #Using Decision tree model
```

```
In [104]: pip install imblearn
```

```
Requirement already satisfied: imblearn in c:\users\asus\anaconda3\lib\site-packages (0.0)
Requirement already satisfied: imbalanced-learn in c:\users\asus\anaconda3\lib\site-packages
(from imblearn) (0.10.1)
Requirement already satisfied: joblib>=1.1.1 in c:\users\asus\anaconda3\lib\site-packages (fr
om imbalanced-learn->imblearn) (1.2.0)
Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\asus\anaconda3\lib\site-packag
es (from imbalanced-learn->imblearn) (1.2.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\asus\anaconda3\lib\site-packa
ges (from imbalanced-learn->imblearn) (2.2.0)
Requirement already satisfied: scipy>=1.3.2 in c:\users\asus\anaconda3\lib\site-packages (fro
m imbalanced-learn->imblearn) (1.7.1)
Requirement already satisfied: numpy>=1.17.3 in c:\users\asus\anaconda3\lib\site-packages (fr
om imbalanced-learn->imblearn) (1.20.3)
Note: you may need to restart the kernel to use updated packages.
```

```
In [105]: from imblearn.over_sampling import SMOTE
        sm = SMOTE()
        X_train_sm, y_train_sm= sm.fit_resample(X_train, y_train.astype('int'))
```

```
In [106]: df2['class'].value_counts()
```

```
Out[106]: 0    128164
          1     13240
          Name: class, dtype: int64
```

```
In [107]: y_train_sm.value_counts()
```

```
Out[107]: 0    102528
          1    102528
          Name: class, dtype: int64
```

```
In [108]: lr=LogisticRegression(random_state=20,solver='liblinear')
lr_model=lr.fit(X_train_sm,y_train_sm)

y_pred_smote=lr_model.predict(X_test)

print(classification_report(y_test,y_pred_smote))
print()
print('RFC score ',np.mean(cross_val_score(lr,X_train_sm,y_train_sm,cv=5)))
print()
print(confusion_matrix(y_test,y_pred_smote))
print()
print('train score ',lr_model.score(X_train_sm,y_train_sm))
print()
print('test score ',lr_model.score(X_test,y_test))
```

	precision	recall	f1-score	support
0	0.97	0.76	0.85	25636
1	0.25	0.78	0.38	2645
accuracy			0.76	28281
macro avg	0.61	0.77	0.61	28281
weighted avg	0.90	0.76	0.81	28281

```
RFC score 0.7516044333180225
```


```
[[19410  6226]
 [  595 2050]]
```

```
train score 0.7516629603620475
```

```
test score 0.7588133375764647
```

```
In [109]: from sklearn.metrics import roc_curve, auc
          from sklearn.model_selection import train_test_split
          from sklearn.metrics import roc_auc_score
          from sklearn.linear_model import LogisticRegression
```

```
In [110]: #create ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_smote)
plt.plot(fpr,tpr)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



Since our model has improved for the Target class we can say that it has improved, but we will try with another models.
Train AND Test are not having equal score which means model is not underfitting.

Using Decision tree model

In []:

```
In [112]: from sklearn.tree import DecisionTreeClassifier
dc=DecisionTreeClassifier(random_state=10)
dc_model=dc.fit(X_train_sm,y_train_sm)

y_pred_dc=dc_model.predict(X_test)

print(classification_report(y_test,y_pred_dc))
print()

print('RFC score ',np.mean(cross_val_score(dc,X_train_sm,y_train_sm,cv=5)))
print()
print(confusion_matrix(y_test,y_pred_dc))
print()
print('train score ',dc_model.score(X_train_sm,y_train_sm))
print()
print('test score ',dc_model.score(X_test,y_test))
```

	precision	recall	f1-score	support
0	0.97	0.92	0.94	25636
1	0.46	0.69	0.56	2645
accuracy			0.90	28281
macro avg	0.72	0.81	0.75	28281
weighted avg	0.92	0.90	0.91	28281


RFC score 0.9164425949212175

```
[[23528 2108]
 [ 813 1832]]
```

train score 1.0

test score 0.8967151090838372

```
In [113]: #create ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_dc)
plt.plot(fpr,tpr)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



#Model is Overfitting

Tuned Random Forest Model

```
In [115]: from sklearn.ensemble import GradientBoostingClassifier, RandomForestClassifier
Rfc=RandomForestClassifier(random_state=10,criterion='gini',max_depth= 5,min_samples_split= 2,
Rfc_modeltune=Rfc.fit(X_train_sm,y_train_sm)

y_pred_tune=Rfc_modeltune.predict(X_test)

print(classification_report(y_test,y_pred_tune))
print()
print('RFC score ',np.mean(cross_val_score(Rfc_modeltune,X_train_sm,y_train_sm,cv=5)))
print()
print(confusion_matrix(y_test,y_pred_tune))
print()
print('train score ',Rfc_modeltune.score(X_train_sm,y_train_sm))
print()
print('test score ',Rfc_modeltune.score(X_test,y_test))
```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	25636
1	0.70	0.70	0.70	2645
accuracy			0.94	28281
macro avg	0.83	0.83	0.83	28281
weighted avg	0.94	0.94	0.94	28281


RFC score 0.8812228159209845

```
[[24838  798]
 [ 796 1849]]
```

train score 0.8812031835205992

test score 0.9436370708249354

```
In [116]: #create ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_tune)
plt.plot(fpr,tpr)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



Using adaboost model

```
In [117]: from sklearn.ensemble import AdaBoostClassifier
abc=AdaBoostClassifier(random_state=10)
abc_model=abc.fit(X_train_sm,y_train_sm)

y_pred_ada=abc_model.predict(X_test)

print(classification_report(y_test,y_pred_ada))
print()

print('RFC score ',np.mean(cross_val_score(abc,X_train_sm,y_train_sm,cv=5)))
print()
print(confusion_matrix(y_test,y_pred_ada))

print('train score ',abc_model.score(X_train_sm,y_train_sm))

print('test score ',abc_model.score(X_test,y_test))
```

	precision	recall	f1-score	support
0	0.97	0.93	0.95	25636
1	0.50	0.71	0.58	2645
accuracy			0.91	28281
macro avg	0.73	0.82	0.77	28281
weighted avg	0.92	0.91	0.91	28281

RFC score 0.9038704119515975

[[23748 1888]

[772 1873]]

train score 0.9067230415106118

test score 0.9059439199462537

In [118]: `pip install parfit==0.220`

```
Requirement already satisfied: parfit==0.220 in c:\users\asus\anaconda3\lib\site-packages (0.220)
Requirement already satisfied: matplotlib in c:\users\asus\anaconda3\lib\site-packages (from parfit==0.220) (3.4.3)
Requirement already satisfied: numpy in c:\users\asus\anaconda3\lib\site-packages (from parfit==0.220) (1.20.3)
Requirement already satisfied: joblib in c:\users\asus\anaconda3\lib\site-packages (from parfit==0.220) (1.2.0)
Requirement already satisfied: sklearn in c:\users\asus\anaconda3\lib\site-packages (from parfit==0.220) (0.0)
Requirement already satisfied: cycycler>=0.10 in c:\users\asus\anaconda3\lib\site-packages (from matplotlib->parfit==0.220) (0.10.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\asus\anaconda3\lib\site-packages (from matplotlib->parfit==0.220) (1.3.1)
Requirement already satisfied: pillow>=6.2.0 in c:\users\asus\anaconda3\lib\site-packages (from matplotlib->parfit==0.220) (8.4.0)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\asus\anaconda3\lib\site-packages (from matplotlib->parfit==0.220) (2.8.2)
Requirement already satisfied: pyparsing>=2.2.1 in c:\users\asus\anaconda3\lib\site-packages (from matplotlib->parfit==0.220) (3.0.4)
Requirement already satisfied: six in c:\users\asus\anaconda3\lib\site-packages (from cycycler->matplotlib->parfit==0.220) (1.16.0)
Requirement already satisfied: scikit-learn in c:\users\asus\anaconda3\lib\site-packages (from sklearn->parfit==0.220) (1.2.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\asus\anaconda3\lib\site-packages (from scikit-learn->sklearn->parfit==0.220) (2.2.0)
Requirement already satisfied: scipy>=1.3.2 in c:\users\asus\anaconda3\lib\site-packages (from scikit-learn->sklearn->parfit==0.220) (1.7.1)
Note: you may need to restart the kernel to use updated packages.
```

In [119]: `from sklearn.model_selection import ParameterGrid`

```
import parfit.parfit as pf

from sklearn.metrics import roc_auc_score, recall_score
#Check first with Adaboost
ad = AdaBoostClassifier()

paramGrid = ParameterGrid({
    'n_estimators': [100,150,175,200,250],
    'learning_rate': [0.1,0.3,0.5,0.7,0.9]
})
```

In []:

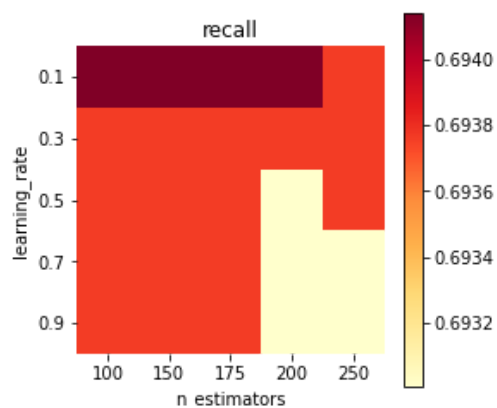
```
In [120]: best_model, best_score, all_models, all_scores = pf.bestFit(ad, paramGrid,
    X_train,y_train,X_test, y_test,
    metric=recall_score, scoreLabel='recall')
```

-----FITTING MODELS-----

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done   5 tasks      | elapsed:   53.2s
[Parallel(n_jobs=-1)]: Done  10 tasks      | elapsed:   1.5min
[Parallel(n_jobs=-1)]: Done  17 tasks      | elapsed:   2.5min
[Parallel(n_jobs=-1)]: Done  21 out of  25 | elapsed:   3.1min remaining:   35.9s
[Parallel(n_jobs=-1)]: Done  25 out of  25 | elapsed:   3.7min finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
```

-----SCORING MODELS-----

```
[Parallel(n_jobs=-1)]: Done  25 out of  25 | elapsed:   19.5s finished
```



```
In [121]: #Tunned Adaboost
```

```
In [122]: from sklearn.ensemble import AdaBoostClassifier
abc=AdaBoostClassifier(random_state=20,n_estimators=100,learning_rate=0.1)
abc_model_tune=abc.fit(X_train_sm,y_train_sm)

y_pred_ada_tune=abc_model.predict(X_test)

print(classification_report(y_test,y_pred_ada_tune))
print()

print('RFC score ',np.mean(cross_val_score(abc,X_train_sm,y_train_sm,cv=5)))
print()
print(confusion_matrix(y_test,y_pred_ada_tune))

print('train score ',abc_model_tune.score(X_train_sm,y_train_sm))

print('test score ',abc_model_tune.score(X_test,y_test))
```

	precision	recall	f1-score	support
0	0.97	0.93	0.95	25636
1	0.50	0.71	0.58	2645
accuracy			0.91	28281
macro avg	0.73	0.82	0.77	28281
weighted avg	0.92	0.91	0.91	28281

RFC score 0.8822224988282183

```
[[23748 1888]
 [ 772 1873]]
```

train score 0.8805106897627965

test score 0.9545277748311587

Gradient Boost Model

```
In [123]: from sklearn.ensemble import GradientBoostingClassifier
gb = GradientBoostingClassifier(random_state=10)
pred_test_gb_sm = gb.fit(X_train_sm,y_train_sm).predict(X_test)

print(classification_report(y_test,pred_test_gb_sm))
print()

print('RFC score ',np.mean(cross_val_score(gb,X_train_sm,y_train_sm,cv=5)))
print()
print(confusion_matrix(y_test,pred_test_gb_sm))

print('train score ',gb.score(X_train_sm,y_train_sm))

print('test score ',gb.score(X_test,y_test))
```

	precision	recall	f1-score	support
0	0.97	0.93	0.95	25636
1	0.51	0.71	0.59	2645
accuracy			0.91	28281
macro avg	0.74	0.82	0.77	28281
weighted avg	0.93	0.91	0.92	28281

RFC score 0.9064550333441357

```
[[23825 1811]
 [ 771 1874]]
train score 0.9067132880774033
test score 0.9087019553764011
```

Using XGboost

```
In [124]: pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\asus\anaconda3\lib\site-packages (1.7.5)
Requirement already satisfied: scipy in c:\users\asus\anaconda3\lib\site-packages (from xgboost) (1.7.1)
Requirement already satisfied: numpy in c:\users\asus\anaconda3\lib\site-packages (from xgboost) (1.20.3)
Note: you may need to restart the kernel to use updated packages.
```

In [125]:

```
from xgboost import XGBClassifier
xgb=XGBClassifier(random_state=10,eval_metric='logloss',use_label_encoder=False)
xgb_model=xgb.fit(X_train_sm,y_train_sm)
y_pred_xgb=xgb_model.predict(X_test)

print(classification_report(y_test,y_pred_xgb))
print()

print('RFC score ',np.mean(cross_val_score(xgb,X_train_sm,y_train_sm,cv=5)))
print()
print(confusion_matrix(y_test,y_pred_xgb))

print('train score ',xgb_model.score(X_train_sm,y_train_sm))

print('test score ',xgb_model.score(X_test,y_test))
```

	precision	recall	f1-score	support
0	0.97	0.94	0.95	25636
1	0.55	0.70	0.62	2645
accuracy			0.92	28281
macro avg	0.76	0.82	0.78	28281
weighted avg	0.93	0.92	0.92	28281

RFC score 0.9127899615573923

```
[[24091 1545]
 [ 781 1864]]
```

train score 0.9221042056803995

test score 0.9177539690958594

Use of Parfit for model tuning

```
In [126]: from sklearn.model_selection import ParameterGrid

import parfit.parfit as pf

from sklearn.metrics import roc_auc_score, recall_score
xg = XGBClassifier()

paramGrid = ParameterGrid({
    'n_estimators': [100,150,175,200,250],
    'learning_rate': [0.1,0.3,0.5,0.7,0.9],
})

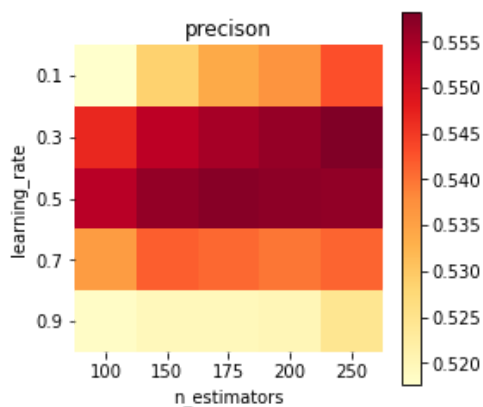
best_model, best_score, all_models, all_scores = pf.bestFit(xg, paramGrid,
    X_train_sm, y_train_sm, X_test, y_test,
    metric=precision_score, scoreLabel='precison')

-----FITTING MODELS-----

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done   5 tasks      | elapsed:   2.0min
[Parallel(n_jobs=-1)]: Done  10 tasks     | elapsed:   3.6min
[Parallel(n_jobs=-1)]: Done  17 tasks     | elapsed:   6.0min
[Parallel(n_jobs=-1)]: Done  21 out of  25 | elapsed:   7.7min remaining:  1.5min
[Parallel(n_jobs=-1)]: Done  25 out of  25 | elapsed:   8.8min finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

-----SCORING MODELS-----

[Parallel(n_jobs=-1)]: Done  25 out of  25 | elapsed:    2.6s finished
```



Tunned XG Boost

```
In [127]: from xgboost import XGBClassifier
xgb=XGBClassifier(random_state=10,eval_metric='logloss',use_label_encoder=False,n_estimators=2000)
xgb_model_tune=xgb.fit(X_train_sm,y_train_sm)
y_pred_xgb_tune=xgb_model.predict(X_test)

print(classification_report(y_test,y_pred_xgb_tune))
print()

print('RFC score ',np.mean(cross_val_score(xgb,X_train_sm,y_train_sm,cv=5)))
print()
print(confusion_matrix(y_test,y_pred_xgb_tune))

print('train score ',xgb_model_tune.score(X_train_sm,y_train_sm))

print('test score ',xgb_model_tune.score(X_test,y_test))
```

	precision	recall	f1-score	support
0	0.97	0.94	0.95	25636
1	0.55	0.70	0.62	2645
accuracy			0.92	28281
macro avg	0.76	0.82	0.78	28281
weighted avg	0.93	0.92	0.92	28281

RFC score 0.9190809251450958

```
[[24091 1545]
 [ 781 1864]]
train score 0.9348373127340824
test score 0.9202291290972738
```

In []:

#CONCLUSIN

We want to pay special attention to accuracy, precision, recall, and the F1 score. Accuracy is a performance metric that is very intuitive: it is simply the ratio of all correctly predicted cases whether positive or negative and all cases in the data.

Precision is the ratio of correctly predicted positive cases vs. all predicted positive cases. It answers the following question: Of all people in the database we predicted as defaulters (fraud), how many were actually fraud? If we think about it, high precision should yield to a low false positive rate. This is important in terms of marketing ,the marketing team has to focus the target customers who ar actually purchasing the products! If we want to maximize our profit, precision is key.

Recall or sensitivity is the ratio of all correctly identified target class divided by all actually target cases. It answers the following question: How well did our model perform in correctly identifying fraud class of customers among all true customers. This is important if we want to maximize overall response (call it revenue maximization)

The F1 Score is the weighted average of precision and recall, hence it takes both false positives and false negatives into account. If data has an uneven class distribution, then the F1 score is far more useful than accuracy. Accuracy works best if false positives and false negatives have similar values. If the value of false positives and false negatives are very different, it's better to look at both precision and recall.

ROC AUC:

ROC stands for Receiver Operating Characteristic, used to evaluate the performance of model . sensitivity is the proportion of correctly predicted events (cases), while specificity is the the proportion of correctly identified non-events (cases). Ideally, both specificity and sensitivity should be high. The ROC curve represents the tradeoff between the two constructs.

The ROC-AUC score provides us with information about how well a model is performing its job of separating cases: in our case distinguishing fraud class from those who are actually making a purchase . Example- 0.92 score means that there is a 92% chance that a model can distinguish fraud from non fraud.

Accuracy: Tuned random forest model best accuracy with score of 94.00 followed by XG Boost Model and tuned XG Boost produced the accuracy with identical scores at 92.00

F1 Score: Tuned Random Forest gave best result

ROC Score: Adaptive Boosting achieved the best score followed by gradient boosting and then XG Boost.

FINAL INFERENCE

Tuned Random Forest model is probably a best choice as it provides the best accuracy and F1 Score It is the best model in separating fraud purchase from non-fraud purchase.

The company would like to segment their customers using various attributes, so that they can perform targeted marketing campaign. Do the following to group the customers in the different clusters

In the original feature engineered dataset, exclude the following columns: ids, dates, class (target column) Convert all the input variables to numeric Using Elbow method identify optimal number of clusters required to group customers in to various clusters Perform K-Means clustering using appropriate number of clusters Tag each customer with a cluster Cluster-wise report average values of all input variables

Using Elbow method identify optimal number of clusters required to group customers in to various clusters

Perform K-Means clustering using appropriate number of clusters Tag each customer with a cluster Cluster-wise report average values of all input variables

```
In [128]: df3 = pd.read_csv('Purchase_Fraud_data.csv')
df3.head(5)
```

Out[128]:

	user_id	signup_time	purchase_time	purchase_value	device_id	source	browser	sex	ip_address	cl
0	22058	2015-02-24 22:55:49	2015-04-18 02:47:11	65278	QVPSPJUOCKZAR	SEO	Chrome	M	7.327584e+08	
1	333320	2015-06-07 20:39:50	2015-06-08 01:38:54	96399	EOGFQPIZPYXFZ	Ads	Chrome	F	3.503114e+08	
2	1359	2015-01-01 18:52:44	2015-01-01 18:52:45	57296	YSSKYOSJHPPLJ	SEO	Opera	M	2.621474e+09	
3	150084	2015-04-28 21:13:25	2015-05-04 13:54:50	43650	ATGTXYKYKUDUQN	SEO	Safari	M	3.840542e+09	
4	221365	2015-07-21 07:09:52	2015-09-09 18:40:53	45016	NAUITBZFJKHWW	Ads	Safari	M	4.155831e+08	

In [129]: df3

Out[129]:

	user_id	signup_time	purchase_time	purchase_value	device_id	source	browser	sex	ip_addres
0	22058	2015-02-24 22:55:49	2015-04-18 02:47:11	65278	QVPSPJUOCKZAR	SEO	Chrome	M	7.327584e+0
1	333320	2015-06-07 20:39:50	2015-06-08 01:38:54	96399	EOGFQPIZPYXFZ	Ads	Chrome	F	3.503114e+0
2	1359	2015-01-01 18:52:44	2015-01-01 18:52:45	57296	YSSKYOSJHPPLJ	SEO	Opera	M	2.621474e+0
3	150084	2015-04-28 21:13:25	2015-05-04 13:54:50	43650	ATGTXYKYKUDUQN	SEO	Safari	M	3.840542e+0
4	221365	2015-07-21 07:09:52	2015-09-09 18:40:53	45016	NAUITBZFJKHWW	Ads	Safari	M	4.155831e+0
...
151107	345170	2015-01-27 03:03:34	2015-03-29 00:30:47	58549	XPSKTWGPWINLR	SEO	Chrome	M	3.451155e+0
151108	274471	2015-05-15 17:43:29	2015-05-26 12:24:39	57952	LYSFABUCPCGBA	SEO	Safari	M	2.439047e+0
151109	368416	2015-03-03 23:07:31	2015-05-20 07:07:47	19003	MEQHCSJUBRBF	SEO	IE	F	2.748471e+0
151110	207709	2015-07-09 20:06:07	2015-09-07 09:34:46	68296	CMCXFGRHYSTVJ	SEO	Chrome	M	3.601175e+0
151111	138208	2015-06-10 07:02:20	2015-07-21 02:03:53	23622	ZINIADFCLHYPG	Direct	IE	M	4.103825e+0

151112 rows × 12 columns

In [130]: df3.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 151112 entries, 0 to 151111
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   user_id         151112 non-null int64
1   signup_time     151112 non-null object
2   purchase_time   151112 non-null object
3   purchase_value  151112 non-null int64
4   device_id       151112 non-null object
5   source          151112 non-null object
6   browser         151112 non-null object
7   sex             146185 non-null object
8   ip_address      151112 non-null float64
9   class           151112 non-null int64
10  category        151112 non-null object
11  dob             146188 non-null object
dtypes: float64(1), int64(3), object(8)
memory usage: 13.8+ MB
```

```
In [131]: import datetime as dt
df3['dob'] = pd.to_datetime(df3['dob'], format = "%d-%m-%Y")
now = pd.Timestamp('now')
df3['age'] = (now - df3['dob'])
df3['age'] = df3['age'].astype(str)
df3[['age', 'age_waste']] = df3['age'].str.split("days", expand=True)
```

```
In [7]: df3 = df3.drop(['age_waste'], axis = 1)
df3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 151112 entries, 0 to 151111
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   user_id                151112 non-null  int64
1   signup_time            151112 non-null  object
2   purchase_time          151112 non-null  object
3   purchase_value         151112 non-null  int64
4   device_id              151112 non-null  object
5   source                 151112 non-null  object
6   browser                151112 non-null  object
7   sex                    146185 non-null  object
8   ip_address             151112 non-null  float64
9   class                  151112 non-null  int64
10  category               151112 non-null  object
11  dob                    146188 non-null  datetime64[ns]
12  age                    151112 non-null  object
dtypes: datetime64[ns](1), float64(1), int64(3), object(8)
memory usage: 15.0+ MB
```

```
In [136]: df3['age'] = pd.to_numeric(df3['age'], errors='coerce')
```

```
In [137]: df3['age'] = df3['age'] / 365
```

```
In [138]: import datetime as dt
df3['tot_mins_diff']=pd.to_datetime(df3['purchase_time'])-pd.to_datetime(df3['signup_time'])

df3.head()

df3['tot_mins_diff'] = df3['tot_mins_diff'].dt.total_seconds()

# Conversion of Difference in Minutes

df3['tot_mins_diff']=df3['tot_mins_diff']/60

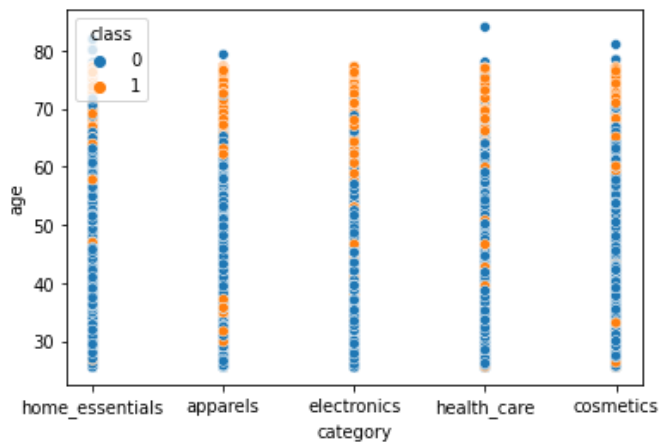
df3.head()
```

Out[138]:

	user_id	signup_time	purchase_time	purchase_value	device_id	source	browser	sex	ip_address	ck
0	22058	2015-02-24 22:55:49	2015-04-18 02:47:11	65278	QVPSPJUOCKZAR	SEO	Chrome	M	7.327584e+08	
1	333320	2015-06-07 20:39:50	2015-06-08 01:38:54	96399	EOGFQIPZYXFZ	Ads	Chrome	F	3.503114e+08	
2	1359	2015-01-01 18:52:44	2015-01-01 18:52:45	57296	YSSKYOSJHPPLJ	SEO	Opera	M	2.621474e+09	
3	150084	2015-04-28 21:13:25	2015-05-04 13:54:50	43650	ATGTXKYKUDUQN	SEO	Safari	M	3.840542e+09	
4	221365	2015-07-21 07:09:52	2015-09-09 18:40:53	45016	NAUITBZFJKHWW	Ads	Safari	M	4.155831e+08	

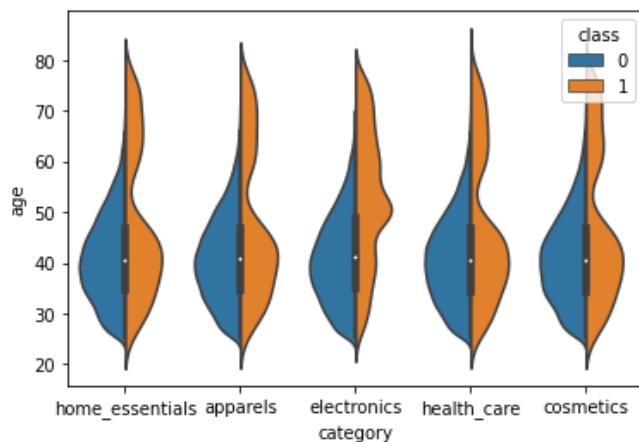
```
In [139]: sns.scatterplot(data=df3, x="category", y="age", hue='class')
```

Out[139]: <AxesSubplot:xlabel='category', ylabel='age'>




```
In [140]: sns.violinplot(data=df3, x="category", y="age", hue="class", split=True)
```

```
Out[140]: <AxesSubplot:xlabel='category', ylabel='age'>
```



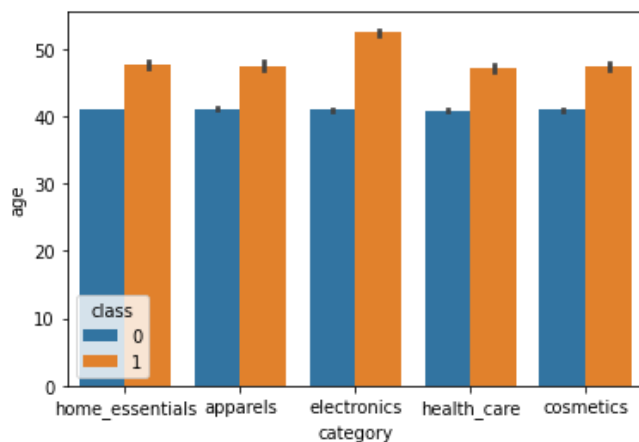
```
In [141]: sns.scatterplot(data=df3, x="age", y="category", hue='class')
```

```
Out[141]: <AxesSubplot:xlabel='age', ylabel='category'>
```



```
In [142]: sns.barplot(data=df3, x="category", y="age", hue="class")
```

```
Out[142]: <AxesSubplot:xlabel='category', ylabel='age'>
```



#We take three features for consideration of creating customer segemntation which are :
'purchase_value','age','tot_mins_diff'

```
In [143]: features=pd.DataFrame(df3, columns = ['purchase_value','age','tot_mins_diff'])
```

```
In [144]: features=pd.get_dummies(features,drop_first=True)
```

```
In [145]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
```

```
In [146]: features.head(5)
```

```
Out[146]:
```

	purchase_value	age	tot_mins_diff
0	65278	47.183562	75111.366667
1	96399	61.331507	299.066667
2	57296	61.161644	0.016667
3	43650	48.824658	8201.416667
4	45016	52.682192	72691.016667

```
In [147]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
```

```
In [148]: features = features.dropna()
```

```
In [149]: features.head(5)
```

```
Out[149]:
```


	purchase_value	age	tot_mins_diff
0	65278	47.183562	75111.366667
1	96399	61.331507	299.066667
2	57296	61.161644	0.016667
3	43650	48.824658	8201.416667
4	45016	52.682192	72691.016667

```
In [150]: #create scaled DataFrame where each variable has mean of 0 and standard dev of 1
scaled_df = StandardScaler().fit_transform(features)
```

```
In [151]: #initialize kmeans parameters
kmeans_kwargs = {
    "init": "random",
    "n_init": 10,
    "random_state": 1,
}

#create list to hold SSE values for each k
sse = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, **kmeans_kwargs)
    kmeans.fit(scaled_df)
    sse.append(kmeans.inertia_)

#visualize results
plt.plot(range(1, 11), sse,color='blue', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)
plt.xticks(range(1, 11))
plt.xlabel("Number of Clusters")
plt.ylabel("SSE")
plt.show()
```



```
#Perform K-Means Clustering with Optimal K=4
```

```
In [153]: #instantiate the k-means class, using optimal number of clusters
kmeans = KMeans(init="random", n_clusters=4, n_init=10, random_state=1)

#fit k-means algorithm to data
kmeans.fit(scaled_df)

#view cluster assignments for each observation
kmeans.labels_
```

```
Out[153]: array([2, 3, 3, ..., 0, 2, 0])
```

```
In [154]: #append cluster assingments to original DataFrame
features['cluster'] = kmeans.labels_

#view updated DataFrame
print(features)
```

	purchase_value	age	tot_mins_diff	cluster
0	65278	47.183562	75111.366667	2
1	96399	61.331507	299.066667	3
2	57296	61.161644	0.016667	3
3	43650	48.824658	8201.416667	3
4	45016	52.682192	72691.016667	3
...
151107	58549	35.720548	87687.216667	2
151108	57952	39.550685	15521.166667	1
151109	19003	33.605479	111360.266667	0
151110	68296	45.238356	85768.650000	2
151111	23622	45.857534	58741.550000	0

[146188 rows x 4 columns]

Getting the total count of the customers belonging to various clusters

```
In [173]: cust1=features[features["cluster"]==0]
print('Number of customer in 1st group=', len(cust1))
```

Number of customer in 1st group= 40388

```
In [174]: cust2=features[features["cluster"]==1]
print('Number of customer in 2nd group=', len(cust2))
```

Number of customer in 2nd group= 41896

```
In [175]: cust3=features[features["cluster"]==2]
print('Number of customer in 3rd group=', len(cust3))
```

Number of customer in 3rd group= 39126

```
In [176]: cust4=features[features["cluster"]==3]
print('Number of customer in 4th group=', len(cust4))
```

Number of customer in 4th group= 24778

```
In [165]: x = features.values
```

```
In [166]: WCSS = []
for i in range(1,11):
    model = KMeans(n_clusters = i,init = 'k-means++')
    model.fit(x)
    WCSS.append(model.inertia_)
fig = plt.figure(figsize = (7,7))
plt.plot(range(1,11),WCSS, linewidth=4, markersize=12,marker='o',color = 'red')
plt.xticks(np.arange(11))
plt.xlabel("Number of clusters")
plt.ylabel("WCSS")
plt.show()
```

C:\Users\asus\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

warnings.warn(

C:\Users\asus\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

warnings.warn(

C:\Users\asus\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

warnings.warn(

C:\Users\asus\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

warnings.warn(

C:\Users\asus\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

warnings.warn(

C:\Users\asus\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

warnings.warn(

C:\Users\asus\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

warnings.warn(

C:\Users\asus\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning


warnings.warn(

C:\Users\asus\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

warnings.warn(


C:\Users\asus\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

warnings.warn(



```
In [167]: model = KMeans(n_clusters = 4, init = "k-means++", max_iter = 300, n_init = 10, random_state =  
y_clusters = model.fit_predict(x)
```

```
In [168]: fig = plt.figure(figsize = (20,20))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x[y_clusters == 0,0],x[y_clusters == 0,1],x[y_clusters == 0,2], s = 10 , color = 'b')
ax.scatter(x[y_clusters == 1,0],x[y_clusters == 1,1],x[y_clusters == 1,2], s = 10 , color = 'o')
ax.scatter(x[y_clusters == 2,0],x[y_clusters == 2,1],x[y_clusters == 2,2], s = 10 , color = 'g')
ax.scatter(x[y_clusters == 3,0],x[y_clusters == 3,1],x[y_clusters == 3,2], s = 10 , color = 'p')
ax.set_xlabel('purchase_value')
ax.set_ylabel('age')
ax.set_zlabel('category')
ax.legend()
plt.show()
```



Customer clusters based on age group , purchase value and tot_mins_diff(difference between sign up time and purchase time)

Customer clusters based on age group , purchase value and tot_mins_diff(difference between sign up time and purchase time
from the above graph we can infer :

cluster 0 indicated with blue has customer base in age group of 40- 80 years where there tot_mins_diff is in range 50000 to 10000 mins and purchase value has high range 0-80000
The total count of customers in cluster 0 is 40388

cluster 1 indicated with orange color has customer base in age group of 30- 80 years where there tot_mins_diff is in range 0 to 25000 mins and purchase value has range 0 to 85000
The total count of customers in cluster 1 is 41896

cluster 2 indicated with green color has customer base in age group of 30- 60 years where there tot_mins_diff is in range 0 to 25000 mins and purchase value has range 0-90000
The total count of customers in cluster 2 is 39126


cluster 3 indicated with purple color has customer base in age group of 40- 90 years where there tot_mins_diff is in range 0 to approximately 80000 mins and purchase value has high range 25000-90000
The total count of customers in cluster 3 is 24778

In [169]: `pip install plotly`

Requirement already satisfied: plotly in c:\users\asus\anaconda3\lib\site-packages (5.14.1)
Requirement already satisfied: tenacity>=6.2.0 in c:\users\asus\anaconda3\lib\site-packages (from plotly) (8.2.2)
Requirement already satisfied: packaging in c:\users\asus\anaconda3\lib\site-packages (from plotly) (21.0)
Requirement already satisfied: pyparsing>=2.0.2 in c:\users\asus\anaconda3\lib\site-packages (from packaging->plotly) (3.0.4)
Note: you may need to restart the kernel to use updated packages.

In [170]: `from mpl_toolkits.mplot3d import Axes3D
import plotly.graph_objs as go
from plotly import tools
from plotly.subplots import make_subplots
import plotly.offline as py`

```
In [177]: Scene = dict(xaxis = dict(title = 'purchase_value'),yaxis = dict(title = 'age'),zaxis = dict(
# model.labels_ is nothing but the predicted clusters i.e y_clusters
labels = model.labels_
trace = go.Scatter3d(x=x[:, 0], y=x[:, 1], z=x[:, 2], mode='markers',marker=dict(color = label
layout = go.Layout(margin=dict(l=0,r=0),scene = Scene,height = 800,width = 800)
data = [trace]
fig = go.Figure(data = data, layout = layout)
fig.show()
```



In []:

